

# Docker, Kubernetes




Motivation, Basics, Hands-on

**Evelyn Haslinger**  
**Markus Zimmermann**

eh@symflower.com  
mz@symflower.com

# Agenda

- **Introduction**
- VMs vs Containers
- Docker
- Kubernetes
- Q&A (... for bigger questions)




# Who Am I?

- Markus Zimmermann <https://github.com/zimmski>
- Writing software since primary school
- Using
  - Docker full-time since 2014
  - Kubernetes full-time since 2016
- **Work** (mainly “project consulting”)
  - Lots of “enterprise” applications, web services, tooling, software infrastructure, distributed apps and clustering, software testing
  - **Now: Software testing and verification @Symflower**

# What is Symflower?

Symflower completely **automatically writes** and runs unit tests revealing bugs, security issues and performance problems.



- Reduce development and maintenance time
- Increase quality of your software and tests

# Example: Find the Problem...

```
func match(s1, s2 []byte) bool {  
    for i := 0; i < len(s1); i++ {  
        c1 := s1[i]  
        c2 := s2[i]  
        if c1 != c2 {  
            c1 |= 'a' - 'A'  
            c2 |= 'a' - 'A'  
            if c1 != c2 || c1 < 'a' || c1 > 'z' {  
                return false  
            }  
        }  
    }  
    return true  
}
```

# Found the Problem!

```
func match(s1, s2 []byte) bool {  
    for i := 0; i < len(s1); i++ {  
        c1 := s1[i]  
        c2 := s2[i] ← Index out of range!  
        if c1 != c2 {  
            c1 |= 'a' - 'A'  
            c2 |= 'a' - 'A'  
            if c1 != c2 || c1 < 'a' || c1 > 'z' {  
                return false  
            }  
        }  
    }  
    return true  
}
```

# Generated Unit Tests

```
func TestMatch117(t *testing.T) {  
    var s1 []byte = []byte{'\x00'}  
    var s2 []byte = nil  
  
    match(s1, s2) // Panic!  
}  
  
func TestMatch119(t *testing.T) {  
    var s1 []byte = []byte{'\x00'}  
    var s2 []byte = []byte{'\x00'}  
  
    actual := match(s1, s2)  
  
    var expected bool = true  
    assert.Equal(t, expected, actual)  
}
```

```
func TestMatch123(t *testing.T) {  
    var s1 []byte = []byte{'c'}  
    var s2 []byte = []byte{'C'}  
  
    actual := match(s1, s2)  
  
    var expected bool = true  
    assert.Equal(t, expected, actual)  
}
```

.....



**Full MC/DC and problem coverage**

If you find that interesting: talk to me or write to [hello@symflower.com](mailto:hello@symflower.com)

# Docker/Kubernetes at Symflower

- We use Docker and Kubernetes **everywhere**
- Services (e.g. issue tracker, mail, websites)
- CI/CD
  - Every pipeline run has at least 3 deployments
  - Every deployment is a Kubernetes deployment
- Symflower the product
  - Every instance is a Kubernetes cluster
  - Every service must be isolated and must scale
  - Every test case is isolated & executed multiple times



# Who are you?

- Who are you?
- What have you done before?
- What are you doing right now?
- **What do you want to achieve today?**

# Agenda

- Introduction
- **VMs vs Containers**
- Docker
- Kubernetes
- Q&A (... for bigger questions)



1

# VMs vs Containers

They seem to be kind of similar

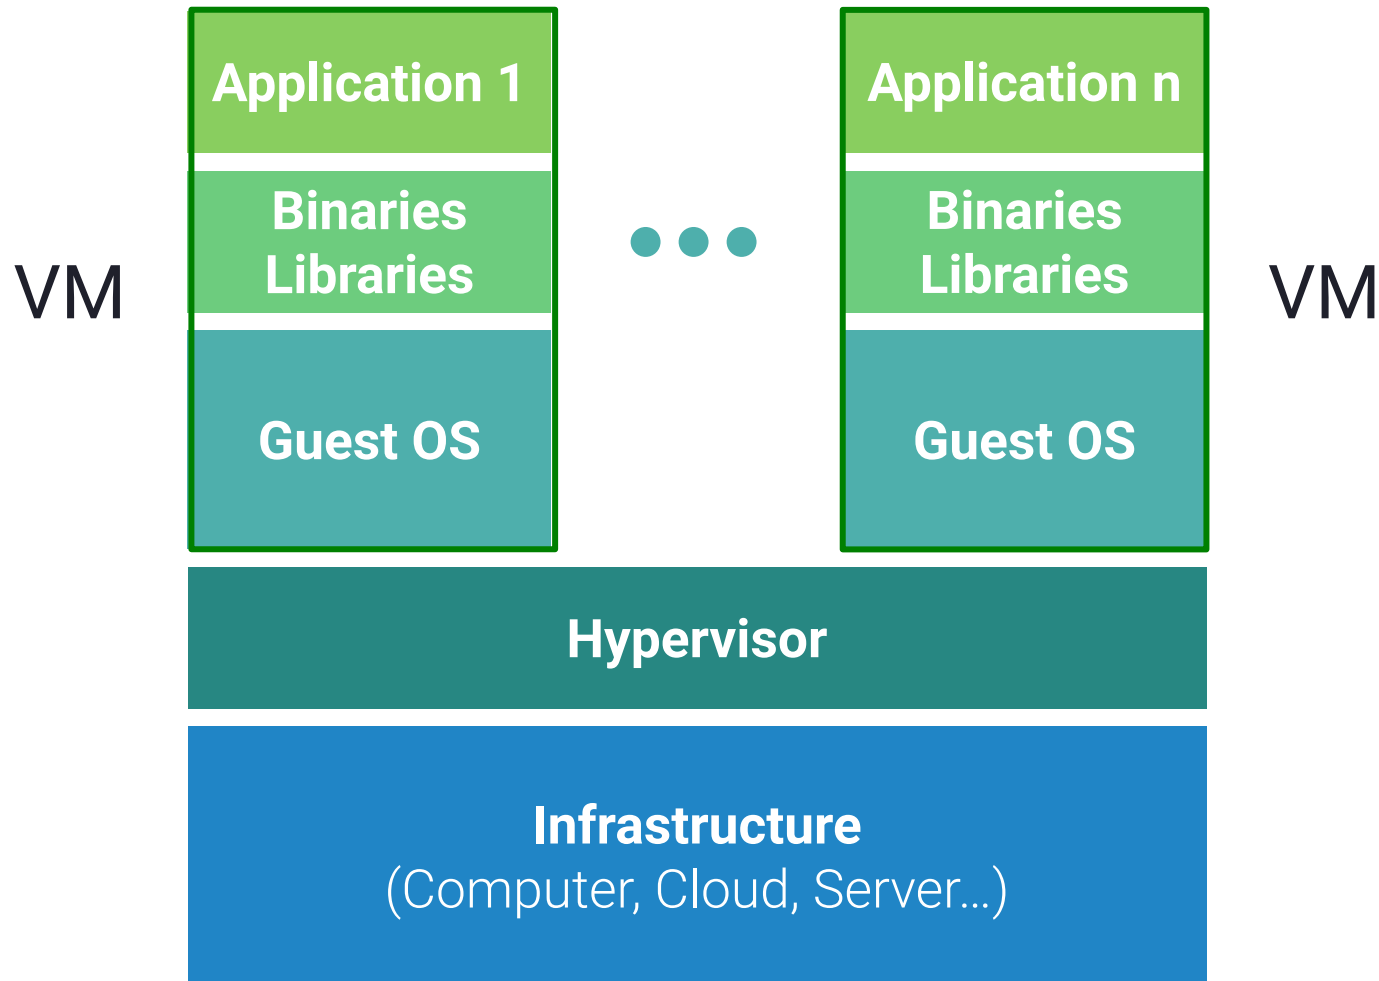
# What is a VM?



A **virtual machine** (VM) is an emulation of a computer system. Virtual machines are based on computer architectures and provide functionality of a physical computer.

[wikipedia.org](https://en.wikipedia.org)

# What is a VM?



# What is a Container?



**OS-level virtualization** refer to an operating system paradigm in which the kernel allows the existence of multiple isolated user-space **instances (e.g. a container)**. Containers may look like real computers from the point of view of programs running in them.

[wikipedia.org](https://wikipedia.org)


# *What is a Container?*



*Think of a container  
as a single process and its children  
that run isolated  
with limited resources.*

Markus Zimmermann

# What is a Container?






# What is a Container?

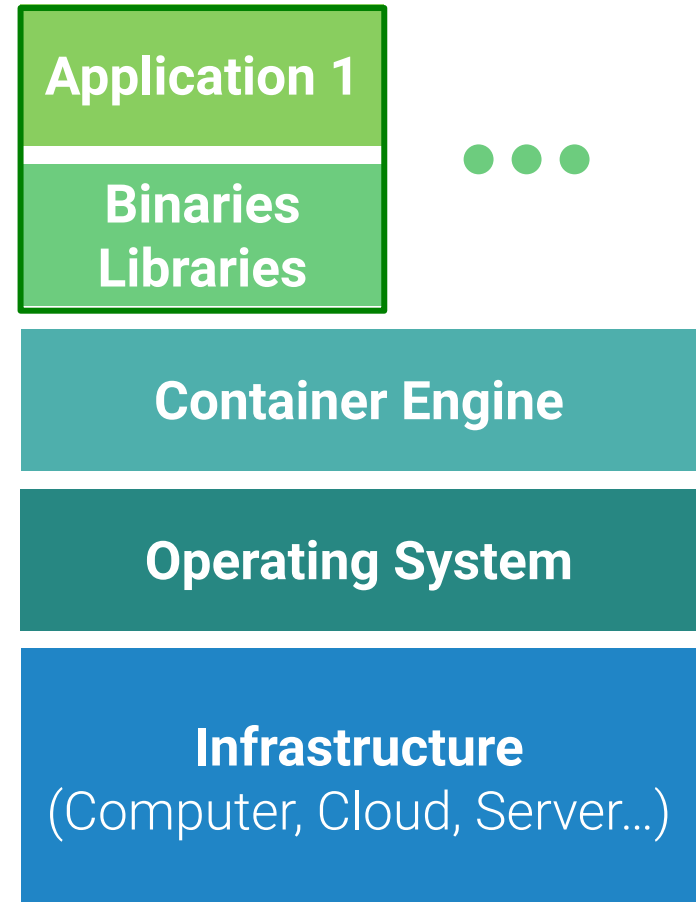
- Nowadays there is a complete ecosystem
- Typical container technology but also ...
- Binaries are replaced with “images”
  - Contain not just files
- Handling of volumes
- Handling of network
- Handling of devices
- Handling of images/life cycle/...

# VMs vs Container?

## VM



## Container



# VMs vs Container?

## VM

- Heavyweight
- Limited performance
- Own OS
- Startup time in minutes
- Fully isolated

## Container

- Lightweight
- Native performance
- Shared OS
- Startup time in seconds
- Process-level isolation

# VMs vs Container?

- When to use one over the other?
  - Use VMs when you do not want to share resources
  - Use VMs when you absolutely do not trust X
  - Use VMs for legacy applications (do not touch a ...)
  - Use containers when you need to maximize the number of applications running on a server
  - Use containers if you need to scale horizontally
  - But in general...

# Should you use containers?

- My honest opinion: **\*\*\*YES\*\*\***
- There are four major reasons
  - Isolation
  - Consistency
  - Reproducibility
  - Utilization
- Using container forces you to think about configurations and to document your decisions

# Btw: Is Container == Docker?

- **No** there are a lot of different technologies
- You can choose e.g.
  - Docker
  - CRI-O
  - Kata Container
  - gVisor
  - ....
- Open Container Initiative (OCI) establishes standards for container technology

# Btw: What about “the Cloud”?

- Do not be fooled: “the cloud” can be on-premise too
- Do not just think about the infrastructure
- It is mainly about modernization, e.g.
  - 12-factor applications
  - Reproducible and reusable
  - DevOps trend
  - “Infrastructure as Code”
  - Rapidly updated applications
  - Efficient resource usage

# Agenda

- Introduction
- VMs vs Containers
- **Docker**
- Kubernetes
- Q&A (... for bigger questions)



# 2

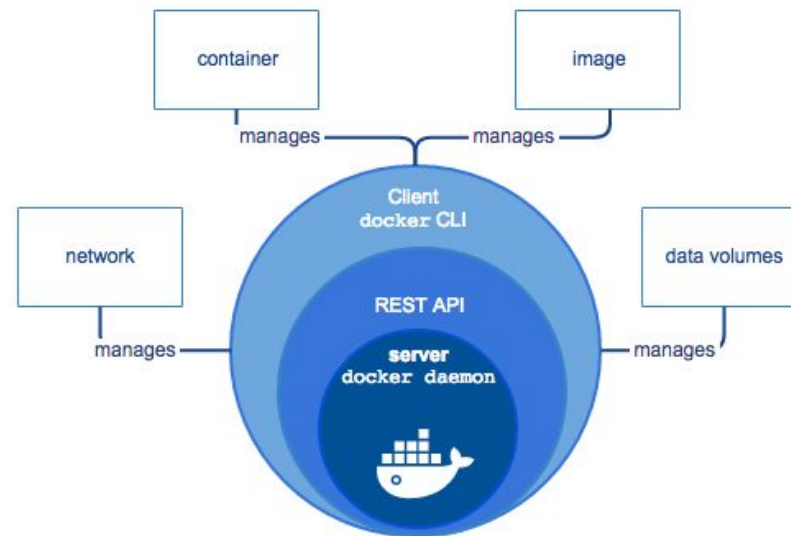
# Docker



The safe and easy way of doing containers

# What is Docker?

- “**Docker** is a set of platform-as-a-service (PaaS) products that use **OS-level virtualization** to deliver software in packages called containers.”



<https://docs.docker.com/engine/docker-overview/>

# Basic Docker Wording 1/2

- Image
  - Basis of a container
  - Basically “the content” e.g. files and defaults
- Registry
  - Stores images
  - Basically an intelligent filesystem for images
- Container
  - One **running** image (not necessarily one service)


# Basic Docker Wording 2/2

- Engine
  - “API” to interact with containers
  - Usually takes care of networking and volumes
- Orchestration (sometimes “control-plane”)
  - Handles life cycle of containers
  - Usually over multiple nodes (hosts)

# What is Moby?

- “Project group” of Docker
- To assemble container based systems
  - And not invent the wheel all over again
- Provides a “Lego-Set” of standard components
  - Orchestration, image management, secret management, networking, provisioning
- Basically:
  - Docker split into Moby
  - Docker is now assembled by Moby projects

# Docker Basics



- The docker file specifies all properties of the container
- Docker files are built into docker images
- Docker images are run in docker containers
- [Docker Hub](#):
  - **THE** common library for container images

# Docker Container Life Cycle

- **Created:** A created container that has not started yet
- **Restarting:** A container currently restarting
- **Started:** A running container
- **Paused:** A container with paused processes
- **Exited:** A container that finished its work
- **Dead:** A container that has been killed

# Docker Namespaces

- Docker namespaces are used for isolation
- Each container has a set of namespaces
- Each aspect of a container runs in a dedicated namespace, its access is limited to this namespace
- E.g. NET, PID, MNT, USER, ...
- Cgroups: used for limiting and isolation resource usage



# Basic Commands

- `docker`
- `docker info`
- `docker build --tag=$tag-name .`
- `docker run $image-name`
- `docker image ls` (or: `docker images`)
- `docker container ls` (or: `docker ps`)

# Agenda

- Introduction
- VMs vs Containers
- Docker
- **Kubernetes**
- Q&A (... for bigger questions)

# 3

# Kubernetes



Infrastructure and deployments made simple



*Kubernetes (k8s) is an **open-source container-orchestration** system for automating deployment, scaling and management of containerized applications.*

[wikipedia.org](https://wikipedia.org)

# Different Perspectives



Developer

The framework for  
deployment and  
infrastructure



Admin

**The OS to  
manage the  
infrastructure**



Manager

Let one person  
perform like ten.

**Standardization + Knowledge of hundreds of experts**


# Fundamental Concepts of K8s 1/2

- Automate everything
  - Resource management
  - Deployments (Rollouts and Rollbacks)
  - Scaling
  - Provisioning
  - Monitoring/Healing/Debugging/...
- Declarative (generic) configuration
  - No explicit host usage
  - No SSH, no scripts -> see Ansible/Puppet/Salt

# Fundamental Concepts of K8s 2/2

- “Infrastructure as Code”
  - Generic configuration and container images
  - Variables define specifics, e.g. user/password of DB
  - Everything is reproducible and reusable
- Everything is disposable (best practice)
  - Pet vs Cattle (see next slide)

# Pet vs. Cattle



- Pets are given names
- They are unique lovingly raised and cared for
- When they get ill, you nurse them back to health



- Cattle is given numbers
- They are identical to one another
- When they get ill, you get another one



# When to (not) use K8s

- Do you have some kind of job/service/server?
  - If not: Sorry, no Kubernetes for you ...
- Single job/service/server
  - Usually a container is fine.
- Multiple ... -> **yes**
- K8s right from the start?
  - Is there k8s experience?
  - Time spent on maintaining
  - ...

# A practical example: CI/CD

- Continuous Integration (CI) is perfect for k8s
  - Goal: Develop, test and release fast
  - Automate builds, checks, deployments
  - Integrate changes multiple times a day
- Continuous Deployment (CD)
  - Automatically deploy each change
- @Symflower
  - Our setup would be almost impossible without k8s
  - Let's take a look

# Installation of K8s

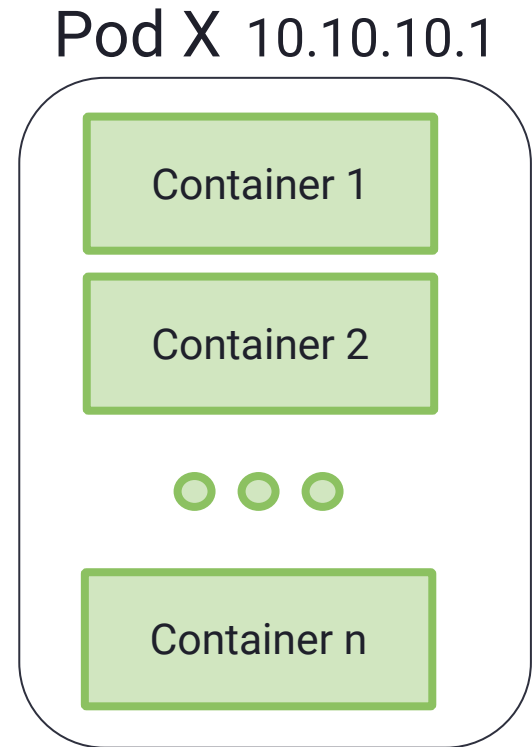
- How and where should I run k8s?
  - [Lots of possibilities](#)
- Your notebook
  - [minikube](#)
  - Vagrant e.g. [with kubespray](#)
- Managed: almost every cloud provider
- On-premise (self-hosted)?
  - WARNING: getting easier every day  
but still knowledge intensive in production

# Basics Concepts of K8s

- We will only look at a minimum
  - **What do you need to deploy an application?**
- What is a k8s ...
  - ... pod
  - ... node
  - ... cluster
  - ... deployment
  - ... service

# K8s Pods

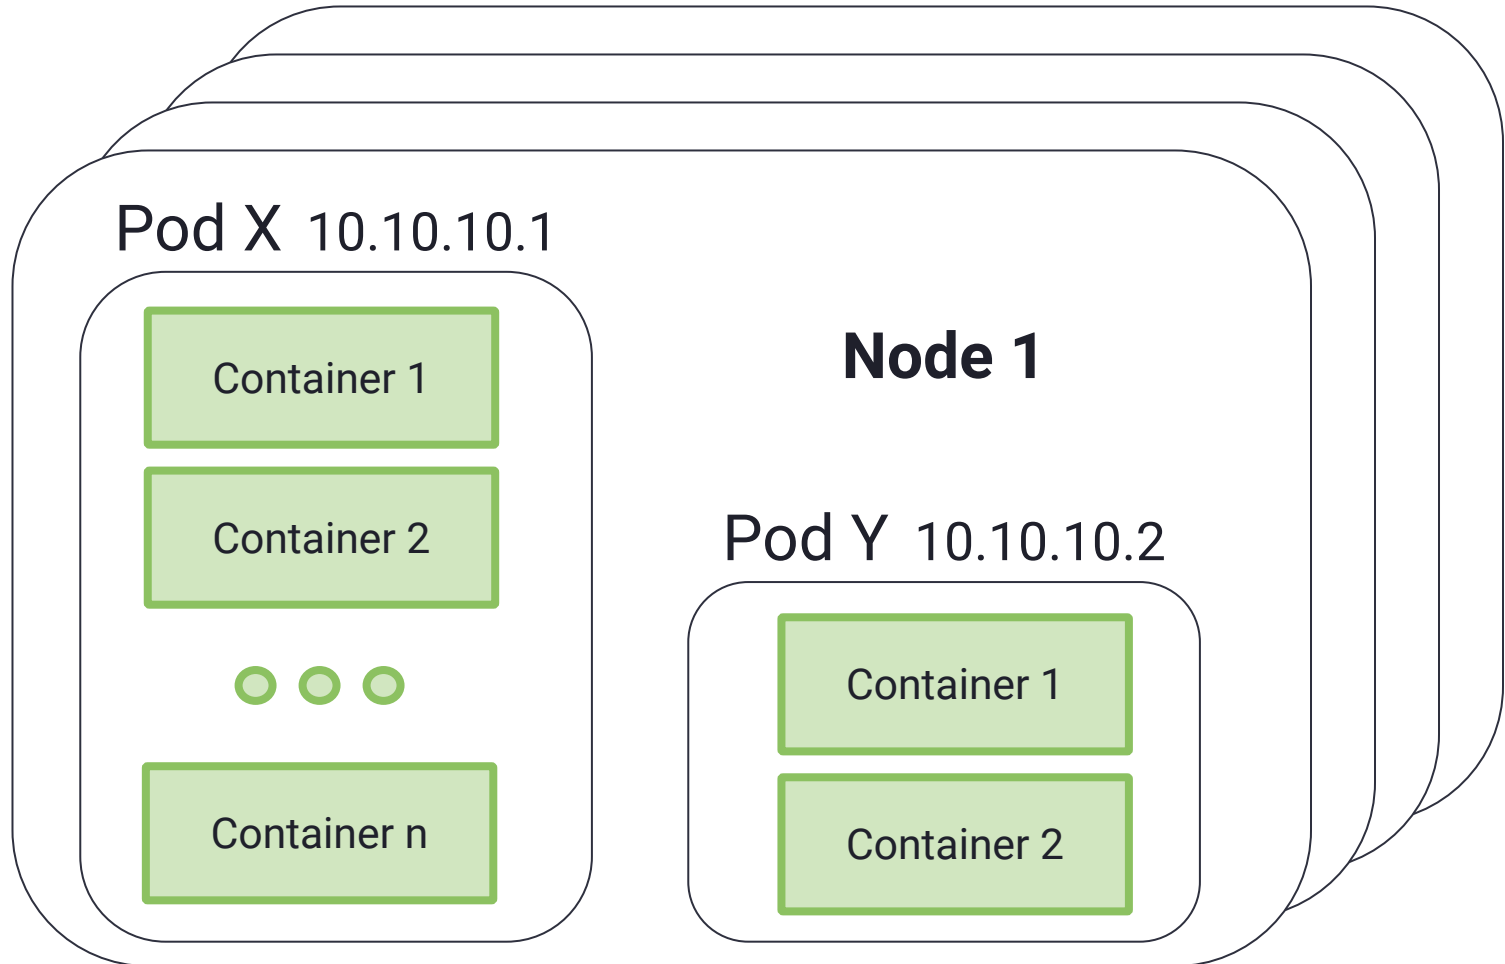
- One or more containers
- Shares resources between its containers
- Needs all of its containers to live
- Why use different containers?
- Which containers should share a pod?



# K8s Nodes and Clusters 1/2

- A node is a “worker machine”
- A cluster consists of one or more nodes
  - Does a single node cluster make sense?
- Nodes provide resources to pods
- A specific pod lives in exactly one node
- Assignment of pods to nodes and their life cycle are managed by k8s
  - What might be good assignment decisions?

# K8s Nodes and Clusters 2/2



# K8s Deployment

- A deployment manages a “replica set” of pods
- **Guarantees the availability of a specific number of identical pods**
- Handles scaling and rollout/rollback of pods






# K8s Service 1/2

- Exposes a set of pods as a network service
  - E.g. deployment “web” with domain “api”
- Not necessarily just one deployment
- Basically an internal load-balancer
  - Pods can be born and die -> routes can change
  - Service is always reachable and forwards

# K8s Service 2/2



# Example: Web Server + Database

- We need all our Docker knowledge but ...
- ... this time we write “infrastructure as code”
- Bonus concept: k8s namespaces
  - A pod has exactly one namespace
  - A namespace is isolated from other namespaces

# Additional K8s Concepts

- ConfigMaps / Secrets
  - Inject configurations
- Persistent Volumes
  - **\*\*Persist\*\*** data to some storage
- Ingress (Controller)
  - Manage external access to services
- Pod/Network/... Policies
  - Restrict access to pods/functionality/...
- ...

# Where do we go from here?

- Lots to learn <https://kubernetes.io/docs/home/>
- Migrating legacy applications?
- Let applications scale?
- Creating a multi-node cluster?
- Creating a cluster of clusters?
- [Manage VMs with k8s?](#)
- [Manage GPUs with k8s?](#)
- ...

# Agenda

- Introduction
- VMs vs Containers
- Docker
- Kubernetes
- **Q&A (... for bigger questions)**

# 4 Q&A





**symflower**  
AUTOMATING QUALITY ASSURANCE

**Evelyn Haslinger**  
**Markus Zimmermann**

eh@symflower.com  
mz@symflower.com



# Feedback

